

Assignment Project Exam Help Parallel Programming

<https://powcoder.com>

Add WeChat powcoder
N-Body Simulation in CUDA

Slides based on Martin Burtscher's tutorial

<https://userweb.cs.txstate.edu/~burtscher/research/ECL-BH/>

Outline

- Review: GPU programming
- N-body example
- Porting and tuning

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



CUDA Programming Model

- Non-graphics programming

- Uses GPU as massively parallel co-processor



- SIMT (single-instruction multiple-threads) model
 - Thousands of threads needed for full efficiency

- C/C++ with extensions

- Function launch
 - Calling functions on GPU
- Memory management
 - GPU memory allocation, copying data to/from GPU
- Declaration qualifiers
 - Device, shared, local, etc.
- Special instructions
 - Barriers, fences, etc.
- Keywords
 - threadIdx, blockIdx

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Calling GPU Kernels

- Kernels are functions that run on the GPU
 - Callable by CPU code
 - CPU can continue processing while GPU runs kernel
- `KernelName<<dim, ...>>>(arg1, arg2, ...);`
- Launch configuration (programmer selectable)
 - GPU spawns m blocks of n threads per block (i.e., $m*n$ threads total) that run a copy of the same function
 - Normal function parameters: passed conventionally
 - Different address space

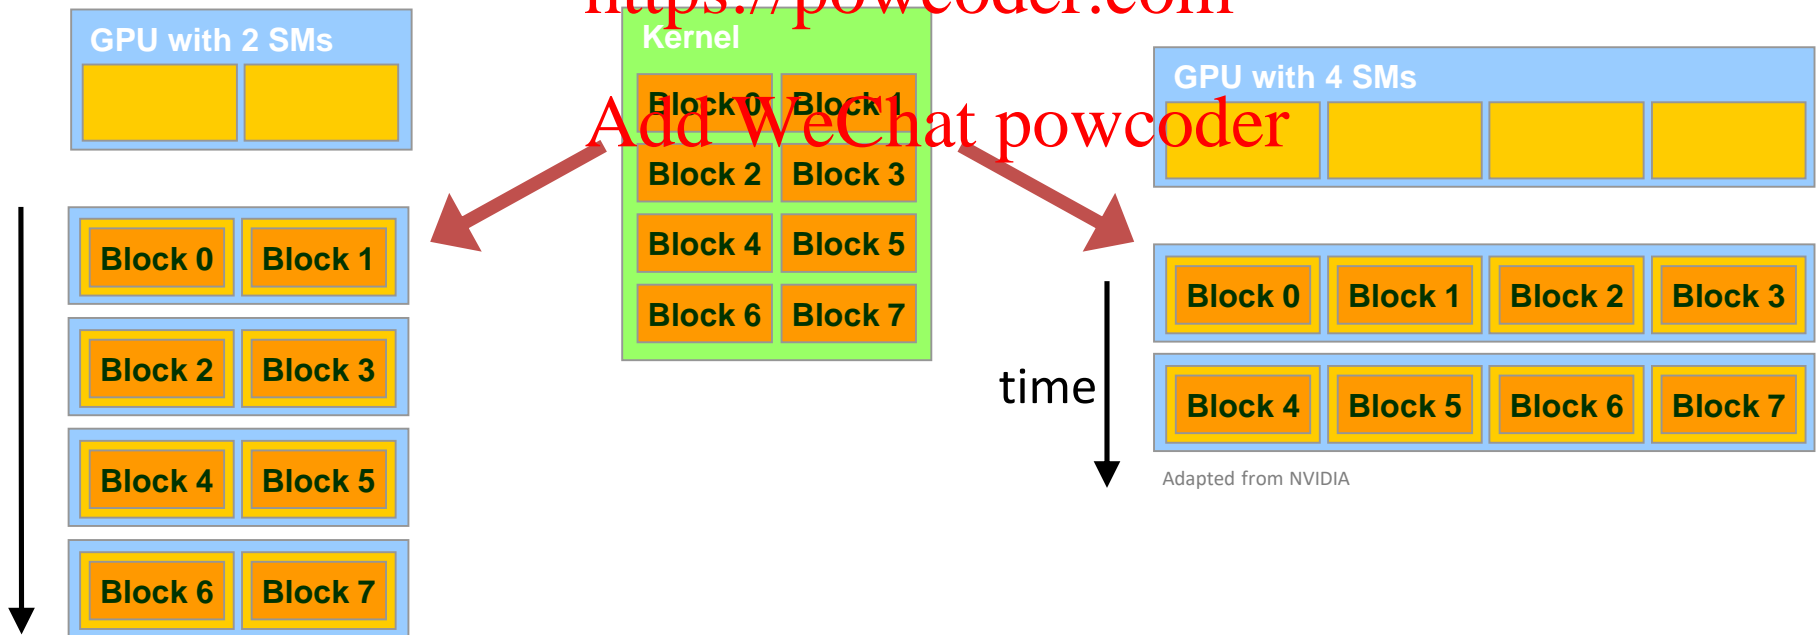
GPU Architecture

- GPUs consist of **Streaming Multiprocessors (SMs)**
 - 1 to 30 SMs per chip (run blocks)
- SMs contain **Processing Elements (PEs)**
 - 8, 32, or 192 PEs per SM (run threads)



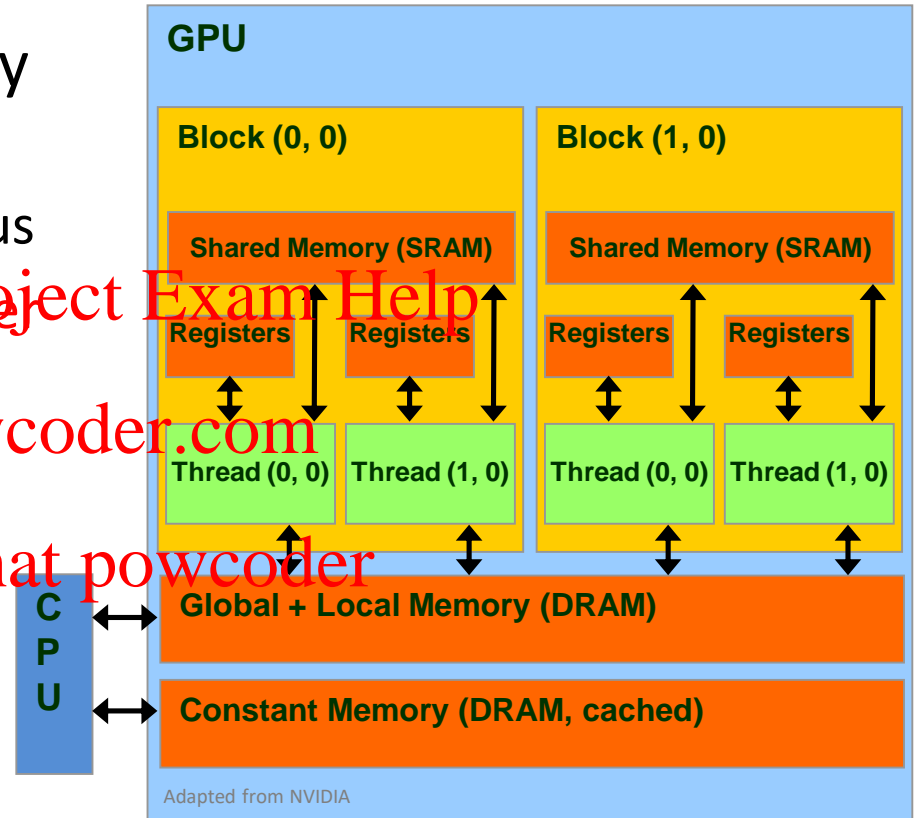
Block Scalability

- Hardware can assign blocks to SMs in any order
 - A kernel with enough blocks scales across GPUs
 - Not all blocks may be resident at the same time



GPU Memories

- Separate from CPU memory
 - CPU can access GPU's global & constant mem. via PCIe bus
 - Requires slow explicit transfer
- Visible GPU memory types
 - Registers (per thread)
 - Local mem. (per thread)
 - Shared mem. (per block)
 - Software-controlled cache
 - Global mem. (per kernel)
 - Constant mem. (read only)



SM Internals (Fermi and Kepler)

- Caches
 - Software-controlled shared memory
 - Hardware-controlled **incoherent** L1 data cache
 - 64 kB combined size, can be split 16/48, 32/32, 48/16
- Synchronization support
 - Fast hardware **barrier** within block (`__syncthreads()`)
 - **Fence** instructions: memory consistency & coherency
- Special operations
 - Thread **voting** (warp-based reduction operations)

Memory Fence Functions

- `void __threadfence_block();`
 - ensures that:
 - All writes to all memory made by the calling thread before the call to `__threadfence_block()` are observed by all threads in the block of the calling thread as occurring before all writes to all memory made by the calling thread after the call to `__threadfence_block();`
 - All reads from all memory made by the calling thread before the call to `__threadfence_block()` are ordered before all reads from all memory made by the calling thread after the call to `__threadfence_block();`

Memory Fence Functions (2)

- `void __threadfence();`
 - acts as `__threadfence_block()` for all threads in the block of the calling thread and also ensures that no writes to all memory made by the calling thread after the call to `__threadfence()` are observed by **any thread in the device** as occurring before any write to all memory made by the calling thread before the call to `__threadfence()`.
 - Note that for this ordering guarantee to be true, the observing threads must truly observe the memory and not cached versions of it; this is ensured by using the `volatile` keyword for the memory variable

Memory Fence Functions (3)

- `void __threadfence_system();`
 - acts as `__threadfence_block()` for all threads in the block of the calling thread and also ensures that all writes to all memory made by the calling thread before the call to `__threadfence_system()` are observed by all threads in the device, host threads, and all threads in peer devices as occurring before all writes to all memory made by the calling thread after the call to `__threadfence_system()`.

Warp Voting Functions

- `int __all_sync(unsigned mask, int predicate);`
 - Evaluate predicate for all non-exited threads in mask and return non-zero if and only if predicate evaluates to non-zero for all of them.
- `int __any_sync(unsigned mask, int predicate);`
- `unsigned __ballot_sync(unsigned mask, int predicate);`
 - Evaluate predicate for all non-exited threads in mask and return an integer whose Nth bit is set if and only if predicate evaluates to non-zero for the Nth thread of the warp and the Nth thread is active.
- `unsigned __activemask();`
 - Returns a 32-bit integer mask of all currently active threads in the calling warp. The Nth bit is set if the Nth lane in the warp is active when `__activemask()` is called.


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Block and Thread Allocation Limits

- Blocks assigned to SMs
 - Until first limit reached
- Hardware limits
 - 8/16 active blocks/SM
 - 1024, 1536, or 2048 resident threads/SM
 - 512 or 1024 threads/blk
 - 16k, 32k, or 64k regs/SM
 - 16 kB or 48 kB shared memory per SM
 - $2^{16}-1$ or $2^{31}-1$ blks/kernel
- Threads assigned to PEs



Warp-based Execution

- 32 contiguous threads form a **warp**
 - Execute **same** instruction in same cycle (or disabled)
 - Warps are scheduled out-of-order with respect to each other to hide latencies
- Thread **divergence**
 - Some threads in warp jump to different PC than others
 - Hardware runs subsets of warp until they re-converge
 - Results in reduction of parallelism (performance loss)


<https://powcoder.com>

Add WeChat powcoder

Thread Divergence

- Non-divergent code

```
if (threadID >= 32) {  
    some_code;  
} else {  
    other_code;  
}
```



- Divergent code

```
if (threadID >= 13) {  
    some_code;  
} else {  
    other_code;  
}
```



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Parallel Memory Accesses

- **Coalesced** main memory access
 - Under some conditions, HW combines multiple (half) warp memory accesses into a single coalesced access
- **Bank-conflict-free** shared memory access
 - No superword alignment or contiguity requirements

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat poweoder

Warnings for GPU Programming

- GPUs can only execute some types of code fast
 - Need lots of data parallelism, data reuse, & regularity
- GPUs are harder to program and tune than CPUs
 - poor tool support
 - architecture
 - poor support for irregular code

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

N-body Simulation


- Time evolution of physical system

- System consists of **bodies**
- “**n**” is the number of bodies
- Bodies interact via **pair-wise forces**

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder



- Many systems can be modeled in this way

- Star/galaxy clusters (gravitational force)
- Particles (electric force, magnetic force)



Simple N-body Algorithm

- Algorithm

Initialize body masses, positions, and velocities

Iterate over time steps {

Accumulate forces acting on each body

Update body positions and velocities based on force

}

Output result

- More sophisticated n-body algorithms exist
 - Barnes Hut algorithm
 - Fast Multipole Method (FMM)

Key Loops (Pseudo Code)

```
bodySet = ...; // input
for timestep do { // sequential
    foreach Body b1 in bodySet { //  $O(n^2)$  parallel
        foreach Body b2 in bodySet {
            if (b1 != b2) {
                b1.addInteractionForce(b2);
            }
        }
    }
    foreach Body b in bodySet { //  $O(n)$  parallel
        b.Advance();
    }
}
// output result
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Force Calculation C Code

```
struct Body {
    float mass, posx, posy, posz; // mass and 3D position
    float velx, vely, velz, accx, accy, accz; // 3D velocity & accel
} *body;

for (i = 0; i < nbodies; i++) {
    ...
    for (j = 0; j < nbodies; j++) {
        if (i != j) {
            dx = body[j].posx - px; // delta x
            dy = body[j].posy - py; // delta y
            dz = body[j].posz - pz; // delta z
            dsq = dx*dx + dy*dy + dz*dz; // distance squared
            dinv = 1.0f / sqrtf(dsq + epssq); // inverse distance
            scale = body[j].mass * dinv * dinv * dinv; // scaled force
            ax += dx * scale; // accumulate x contribution of accel
            ay += dy * scale; az += dz * scale; // ditto for y and z
        }
    }
    ...
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

N-body Algorithm Suitability for GPU

- Lots of data parallelism
 - Force calculations are independent
 - Should be able to keep SMs and PEs busy
- Sufficient memory access regularity
 - All force calculations access body data in same order
 - Should have lots of coalesced memory accesses
- Sufficient code regularity
 - All force calculations are identical
 - There should be little thread divergence
- Plenty of data reuse
 - $O(n^2)$ operations on $O(n)$ data
 - CPU/GPU transfer time is insignificant

C to CUDA Conversion

- Two CUDA kernels

- Force calculation
- Advance position and velocity

Assignment Project Exam Help

<https://powcoder.com>

- Benefits

- Force calculation requires over 99.9% of runtime
 - Primary target for acceleration
- Advancing kernel unimportant to runtime
 - But allows to keep data on GPU during entire simulation
 - Minimizes GPU/CPU transfers

Add WeChat powcoder

C to CUDA Conversion

```
__global__ void ForceCalcKernel(int nbodies, struct Body *body, ...) {  
    . . .  
}  
__global__ void AdvancingKernel(int nbodies, struct Body *body, ...) {  
    . . .  
}  
  
int main(...) {  
    Body *body, *body1;  
    . . .  
    cudaMalloc((void**)&body1, sizeof(Body)*nbodies);  
    cudaMemcpy(body1, body, sizeof(Body)*nbodies, cudaMemcpyHostToDevice);  
    for (timestep = ...) {  
        ForceCalcKernel<<<1, 1>>>(nbodies, body1, ...);  
        AdvancingKernel<<<1, 1>>>(nbodies, body1, ...);  
    }  
    cudaMemcpy(body, body1, sizeof(Body)*nbodies, cudaMemcpyDeviceToHost);  
    cudaFree(body1);  
    . . .  
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Evaluation Methodology

- Systems and compilers
 - CC 1.3: Quadro FX 5800, nvcc 3.2
 - 30 SMs, 240 PEs, 1.3 GHz, 30720 resident threads
 - CC 2.0: Tesla C2050, nvcc 3.2
 - 14 SMs, 448 PEs, 1.15 GHz, 21504 resident threads
 - CC 3.0: GeForce GTX 680, nvcc 4.2
 - 8 SMs, 1536 PEs, 1.0 GHz, 16384 resident threads
- Input and metric
 - 1k, 10k, or 100k star clusters (Plummer model)
 - Median runtime of three experiments, excluding I/O

1-Thread Performance

- Problem size
 - n=10000, step=1
 - n=10000, step=1
 - n=3000, step=1
- Performance
 - 1 thread is one to two orders of magnitude slower on GPU than CPU
- Slowdown rel. to CPU
 - CC 1.3: 72.4
 - CC 2.0: 36.7
 - CC 3.0: 68.1

(Note: comparing different GPUs to different CPUs)
- Reasons
 - No caches (CC 1.3)
 - Not superscalar
 - Slower clock frequency
 - No SMT latency hiding

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Using N Threads

- Approach
 - Eliminate outer loop
 - Instantiate ~~n~~ copies of inner loop, one per body
- Threading
 - Blocks can only hold 512 or 1024 threads
 - Up to 8/16 blocks can be resident in an SM at a time
 - SM can hold 1024, 1536, or 2048 threads
 - Use 256 threads per block (works for all three GPUs)
 - Need multiple blocks
 - Last block may not need all of its threads

Using N Threads

```
__global__ void ForceCalcKernel(int nbodies, struct Body *body, ...) {  
  for (i = 0; i < nbodies; i++) {  
    i = threadIdx.x + blockIdx.x * blockDim.x; // compute i  
    if (i < nbodies) { // in case last block is only partially used  
      for (j = ...) {  
        . . .  
      }  
    }  
  }  
}  
  
__global__ void AdvancingKernel(int nbodies,...) // same changes  
  
#define threads 256  
int main(...) {  
  . . .  
  int blocks = (nbodies + threads - 1) / threads; // compute block cnt  
  for (timestep = ...) {  
    ForceCalcKernel<<<1, 1blocks, threads>>>(nbodies, body1, ...);  
    AdvancingKernel<<<1, 1blocks, threads>>>(nbodies, body1, ...);  
  }  
}
```

Assignment Project Exam Help

<https://powcoder.com>

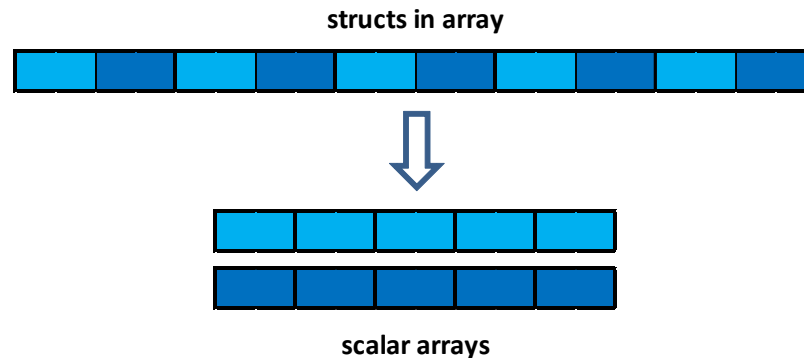
Add WeChat powcoder

N Thread Speedup

- Relative to 1 GPU thread
 - CC 1.3: 7781 (240 PEs)
 - CC 2.0: 6495 (448 PEs)
 - CC 3.0: 12150 (1536 PEs)
- Performance
 - Speedup much higher than number of PEs (32, 14.5, and 7.9 times)
 - Due to SMT latency hiding
- Relative to 1 CPU thread
 - CC 1.3: 107.5
 - CC 2.0: 176.7
 - CC 3.0: 176.2
- Per-core performance
 - CPU core delivers up to 4.4, 5, and 8.7 times as much performance as a GPU core (PE)

Using Scalar Arrays

- Data structure conversion
 - Arrays of structs are bad for coalescing
 - Bodies' elements (e.g., mass fields) are not adjacent
- Optimize data structure
 - Use multiple scalar arrays, one per field (need 10)
 - Results in code bloat but often much better speed



Using Scalar Arrays

```
__global__ void ForceCalcKernel(int nbodies, float *mass, ...) {  
    // change all "body[k].blah" to "blah[k]"  
}  
  
__global__ void AdvancingKernel(int nbodies, float *mass, ...) {  
    // change all "body[k].blah" to "blah[k]"  
}  
  
int main(...) {  
    float *mass, *posx, *posy, *posz, *velx, *vely, *velz, *accx, *accy, *accz;  
    float *mass1, *posx1, *posy1, *posz1, *velx1, *vely1, *velz1, ...;  
    mass = (float *)malloc(sizeof(float) * nbodies); // etc  
    . . .  
    cudaMalloc((void**)&mass1, sizeof(float)*nbodies); // etc  
    cudaMemcpy(mass1, mass, sizeof(float)*nbodies, cuda...HostToDevice); // etc  
    for (timestep = ...) {  
        ForceCalcKernel<<<blocks, threads>>>(nbodies, mass1, posx1, ...);  
        AdvancingKernel<<<blocks, threads>>>(nbodies, mass1, posx1, ...);  
    }  
    cudaMemcpy(mass, mass1, sizeof(float)*nbodies, cuda...DeviceToHost); // etc  
    . . .  
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Scalar Array Speedup

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- Performance
 - Threads access **same** memory locations, not adjacent ones
 - Always combined but not really coalesced access
 - Slowdowns may be due to DRAM page/TLB misses
- Relative to struct
 - CC 1.3: **0.83**
 - CC 2.0: **0.96**
 - CC 3.0: **0.82**
- Scalar arrays
 - Still needed (see later)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Constant Kernel Parameters

- Kernel parameters
 - Lots of parameters due to scalar arrays
 - All but one parameter never change their value
- Constant memory
 - “Pass” parameters only once
 - Copy them into GPU’s constant memory
- Performance implications
 - Reduced parameter passing overhead
 - Constant memory has hardware cache

Constant Kernel Parameters

```
__constant__ int nbodiesd;
__constant__ float dthfd, epssqd, float *massd, *posxd, ...;

__global__ void ForceCalcKernel(int step) {
    // rename affected variables (add "d" to name)
}

__global__ void AdvancingKernel() {
    // rename affected variables (add "d" to name)
}

int main(...) {
    . . .
    cudaMemcpyToSymbol(massd, &massl, sizeof(void *)); // etc
    . . .
    for (timestep = ...) {
        ForceCalcKernel<<<1, 1>>>(step);
        AdvancingKernel<<<1, 1>>>();
    }
    . . .
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Constant Mem. Parameter Speedup

- Problem size
 - n=1000, step=10000
 - n=1000, step=10000
 - n=3000, step=10000
- Performance
 - Minimal perf. impact
 - May be useful for very short kernels that are often invoked
- Speedup
 - CC 1.3: 1.015
 - CC 2.0: 1.016
 - CC 3.0: 0.971
- Benefit
 - Less shared memory used on CC 1.3 devices

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Using the RSQRTF Instruction

- Slowest kernel operation
 - Computing one over the square root is very slow
 - GPU has slightly imprecise but fast $1/\sqrt{}$ instruction
- IEEE floating-point accuracy compliance
 - CC 1.x is not entirely compliant
 - CC 2.x and above are compliant but also offer faster non-compliant instructions

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Using the RSQRT Instruction

```
for (i = 0; i < nbodies; i++) {  
    . . .  
    for (j = 0; j < nbodies; j++) {  
        if (i != j) {  
            dx = body[j].posx - px;  
            dy = body[j].posy - py;  
            dz = body[j].posz - pz;  
            dsq = dx*dx + dy*dy + dz*dz;  
            dinv = 1.0f / sqrtf(dsq + epssq);  
            dinv = rsqrtf(dsq + epssq);  
            scale = body[j].mass * dinv * dinv * dinv;  
            ax += dx * scale;  
            ay += dy * scale;  
            az += dz * scale;  
        }  
    }  
    . . .  
}
```

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

RSQRT Speedup

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- Performance
 - Little change for CC 1.3
 - Compiler automatically uses less precise RSQRTF as most FP ops are not fully precise anyhow
 - 83% speedup for CC 2.0
 - Over entire application
 - Compiler defaults to precise instructions
 - Explicit use of RSQRTF indicates imprecision okay
- Speedup
 - CC 1.3: 0.99
 - CC 2.0: 1.83
 - CC 3.0: 1.64

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Using 2 Loops to Avoid If Statement

- “if (i != j)” creates code divergence
 - Break loop into two loops to avoid if statement

Assignment Project Exam Help

```
for (j = 0; j < nbodies; j++) {  
    if (i != j) {  
        dx = body[j].posx - px;  
        dy = body[j].posy - py;  
        dz = body[j].posz - pz;  
        dsq = dx*dx + dy*dy + dz*dz;  
        dinv = rsqrtf(dsq + epssq);  
        scale = body[j].mass * dinv * dinv * dinv;  
        ax += dx * scale;  
        ay += dy * scale;  
        az += dz * scale;  
    }  
}
```

<https://powcoder.com>

Add WeChat powcoder

Using 2 Loops to Avoid If Statement

```
for (j = 0; j < i; j++) {
    dx = body[j].posx - px;
    dy = body[j].posy - py;
    dz = body[j].posz - pz;
    dsq = dx*dx + dy*dy + dz*dz;
    dinv = rsqrtf(dsq + epssq);
    scale = body[j].mass * dinv * dinv * dinv;
    ax += dx * scale;
    ay += dy * scale;
    az += dz * scale;
}
for (j = i+1; j < bodies; j++) {
    dx = body[j].posx - px;
    dy = body[j].posy - py;
    dz = body[j].posz - pz;
    dsq = dx*dx + dy*dy + dz*dz;
    dinv = rsqrtf(dsq + epssq);
    scale = body[j].mass * dinv * dinv * dinv;
    ax += dx * scale;
    ay += dy * scale;
    az += dz * scale;
}
```

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

Loop Duplication Speedup

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- Performance
 - No change for 2.0 & 3.0
 - Divergence moved to loop
 - 45% slowdown for CC 1.3
 - Unclear reason
- Speedup
 - CC 1.3: 0.55
 - CC 2.0: 1.00
 - CC 3.0: 1.00
- Discussion
 - Not a useful optimization
 - Code bloat
 - A little divergence is okay (only 1 in 3125 iterations)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Blocking using Shared Memory

- Code is memory bound
 - Each warp streams in all bodies' masses and positions
- Use shared memory in inner loop
 - Read block of mass & position info into shared mem
 - Requires barriers (fast hardware barrier within SM)
- Advantage
 - A lot fewer main memory accesses
 - Remaining main memory accesses are fully coalesced (due to usage of scalar arrays)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Blocking using Shared Memory

```
__shared__ float posxs[threads], posys[threads], poszs[...], masss[...];
j = 0;
for (j1 = 0; j1 < nbodiesd; j1 += THREADS) { // first part of loop
    idx = tid + j1;
    if (idx < nbodiesd) { // each thread copies 4 words (fully coalesced)
        posxs[id] = posxd[idx]; posys[id] = posyd[idx];
        poszs[id] = poszd[idx]; masss[id] = massd[idx];
    }
    __syncthreads(); // wait for all copying to be done
    bound = min(nbodiesd - j1, THREADS);
    for (j2 = 0; j2 < bound; j2++) { // second part of loop
        if (i != j) {
            dx = posxs[j2] - px; dy = posys[j2] - py; dz = poszs[j2] - pz;
            dsq = dx*dx + dy*dy + dz*dz;
            dinv = rsqrtf(dsq + epssqd);
            scale = masss[j2] * dinv * dinv * dinv;
            ax += dx * scale; ay += dy * scale; az += dz * scale;
        }
    }
    __syncthreads(); // wait for all force calculations to be done
}
```

Blocking Speedup

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- Performance
 - Great speedup for CC 1.3
 - Some speedup for others
 - Has hardware data cache
- Speedup
 - CC 1.3: 3.7
 - CC 2.0: 1.1
 - CC 3.0: 1.6
- Discussion
 - Very important optimization for memory bound code
 - Even with L1 cache

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Loop Unrolling

- CUDA compiler
 - Generally good at unrolling loops with fixed bounds
 - Does not unroll inner loop of our example code
- Use pragma to unroll (and pad arrays)

```
#pragma unroll 8
for (j2 = 0; j2 < bound; j2++, j++) {
    if (i != j) {
        dx = posxs[j2] - px;  dy = posys[j2] - py;  dz = poszs[j2] - pz;
        dsq = dx*dx + dy*dy + dz*dz;
        dinv = rsqrtf(dsq + epssqd);
        scale = masss[j2] * dinv * dinv * dinv;
        ax += dx * scale;  ay += dy * scale;  az += dz * scale;
    }
}
```

Loop Unrolling Speedup

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- Performance
 - Insignificant speedup
 - All three GPUs
- Discussion
 - Can be useful
 - May increase register usage, which may lower maximum number of threads per block and result in slowdown
- Speedup
 - CC 1.3: 1.07
 - CC 2.0: 1.16
 - CC 3.0: 1.07

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

CC 2.0 Absolute Performance

- Problem size
 - $n=100000$, $\text{step}=1$
- Runtime
 - 612 ms
- FP operations
 - 326.7 GFlop/s
- Main mem throughput
 - 1.035 GB/s
- Not peak performance
 - Only 32% of 1030 GFlop/s
 - Peak assumes FMA every cycle
 - 3 sub (1c), 3 fma (1c), 1 rsqrt (8c), 3 mul (1c), 3 fma (1c) = 20c for 20 Flop
 - 63% of realistic peak of 515.2 GFlop/s
 - Assumes no non-FP operations
 - With int ops = 31c for 20 Flop
 - 99% of actual peak of 330.45 GFlop/s

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Eliminating the If Statement

- Algorithmic optimization
 - Potential softening parameter avoids division by zero
 - If-statement is not necessary and can be removed
 - Eliminates thread divergence

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```
for (j2 = 0; j2 < bound; j2++, j++) {  
  if (i != j) {  
    dx = posxs[j2] - px; dy = posys[j2] - py; dz = poszs[j2] - pz;  
    dsq = dx*dx + dy*dy + dz*dz;  
    dinv = rsqrtf(dsq + epssqd);  
    scale = masss[j2] * dinv * dinv * dinv;  
    ax += dx * scale; ay += dy * scale; az += dz * scale;  
  }  
}
```


If Elimination Speedup

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- Performance
 - Large speedup
 - All three GPUs
- Discussion
 - No thread divergence
 - Allows compiler to schedule code much better
- Speedup
 - CC 1.3: 1.38
 - CC 2.0: 1.54
 - CC 3.0: 1.64

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Rearranging Terms

- Generated code is suboptimal
 - Compiler does not emit as many fused multiply-add (FMA) instructions as it could
 - Rearrange terms in expressions to help compiler
 - Need to check generated assembly code

```
for (j2 = 0; j2 < bound; j2++, j++) {  
    dx = posxs[j2] - px;  dy = posys[j2] - py;  dz = poszs[j2] - pz;  
    dsq = dx*dx + dy*dy + dz*dz;  
    dinv = rsqrtf(dsq + epssqd);  
    dsq = dx*dx + (dy*dy + (dz*dz + epssqd));  
    dinv = rsqrtf(dsq);  
    scale = masss[j2] * dinv * dinv * dinv;  
    ax += dx * scale;  ay += dy * scale;  az += dz * scale;  
}
```

FMA Speedup

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- Performance
 - Small speedup
 - All three GPUs
- Discussion
 - Seemingly needless transformations may make a difference
- Speedup
 - CC 1.3: 1.03
 - CC 2.0: 1.05
 - CC 3.0: 1.06

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Higher Unroll Factor

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- Unroll 128 times
 - Avoid looping overhead
 - Now that there are no *ifs*
- Performance
 - Little speedup/slowdown
- Speedup
 - CC 1.3: 1.01
 - CC 2.0: 1.04
 - CC 3.0: 0.93
- Discussion
 - Carefully choose unroll factor (manually tune)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Compiler Flags

- Problem size
 - n=100000, step=1
 - n=100000, step=1
 - n=300000, step=1
- -use_fast_math
 - “-ftz=true” suffices (flush denormals to zero)
 - Makes SP-FP operations faster except on CC 1.3
- Speedup
 - CC 1.3: 1.00
 - CC 2.0: 1.18
 - CC 3.0: 1.15
- Performance
 - Significant speedup
- Discussion
 - Use faster but less precise operations when prudent

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Final Absolute Performance

- CC 2.0 Fermi GTX 480
 - Problem size
 - $n=100000$, step=1
 - Runtime
 - 296.1 ms
 - FP operations
 - 675.6 GFlop/s (SP)
 - 66% of peak performance
 - 261.1 GFlops/s (DP)
 - Main mem throughput
 - 2.139 GB/s
- CC 3.0 Kepler GTX 680
 - Problem size
 - $n=300000$, step=1
 - Runtime
 - 1073 ms
 - FP operations
 - 1677.6 GFlop/s (SP)
 - 54% of peak performance
 - 88.7 GFlops/s (DP)
 - Main mem throughput
 - 5.266 GB/s

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Hybrid Execution

- CPU always needed for program launch and I/O
 - CPU much faster on serial program segments
- GPU 10 times faster than CPU on parallel code
 - Running 10% of problem on CPU is hardly worthwhile
 - Complicates programming and requires data transfer
 - Best CPU data structure is often not best for GPU
- PCIe bandwidth much lower than GPU bandwidth
 - 1.6 to 6.5 GB/s versus 192 GB/s
 - But can send data while CPU & GPU are computing
 - Merging CPU and GPU on same die (e.g., AMD's Fusion APU) makes finer grain switching possible

Summary

- Step-by-step porting and tuning of CUDA code
 - Example: n-body simulation
- GPUs have very powerful hardware
 - But only exploitable with some code
 - Harder to program and optimize than CPU hardware

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder